



Accelerating OpenCV in Embedded Vision

Stephen Neuendorffer

Principal Engineer, High Level Synthesis

Xilinx

#TheVisionShow

© Copyright 2014 Xilinx

- Trends in vision processing
 - ?
- Key bottlenecks in embedded video processing
 - ?
- Using OpenCV in embedded video processing
 - ?

Increasing Video Resolutions

Time

2013

4K

Video Format

2005

1080p

2000

2K

1995

720p

DVD

VCD

Pixel rate @ 60 fps
M Pix/s

480

120

20



Studio / Cinema Camera



Machine Vision



Video Conferencing



HD Surveillance



Driver Assist

High Resolution Video moving to Camera/Vision applications

Multiple Cameras



Machine Vision



Video Conferencing



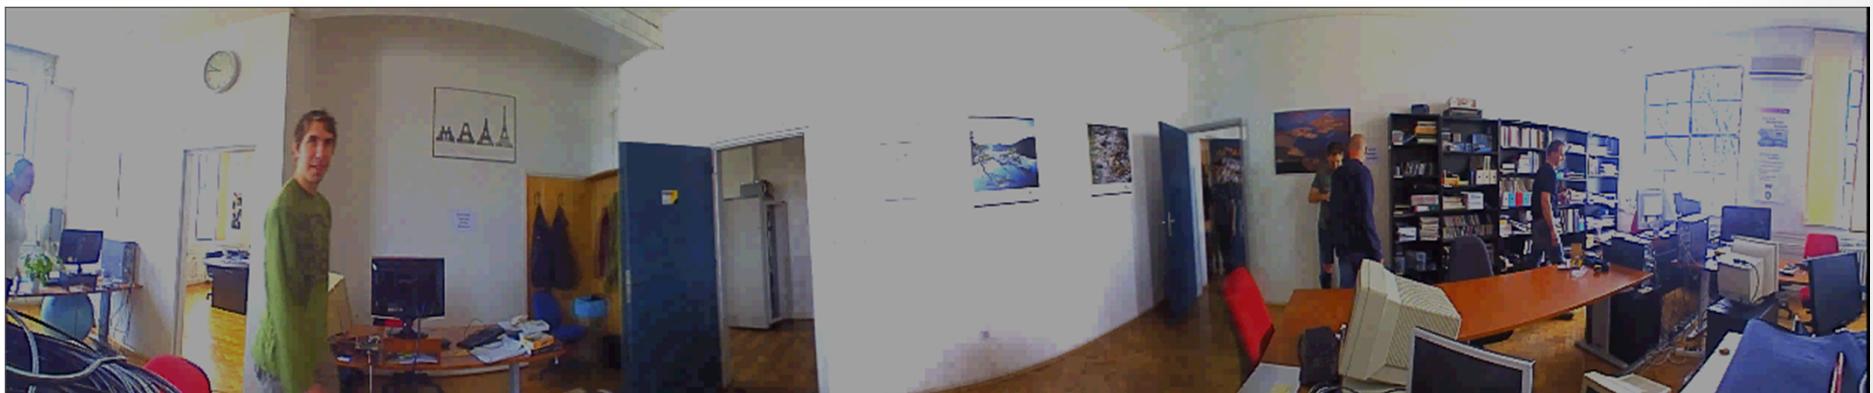
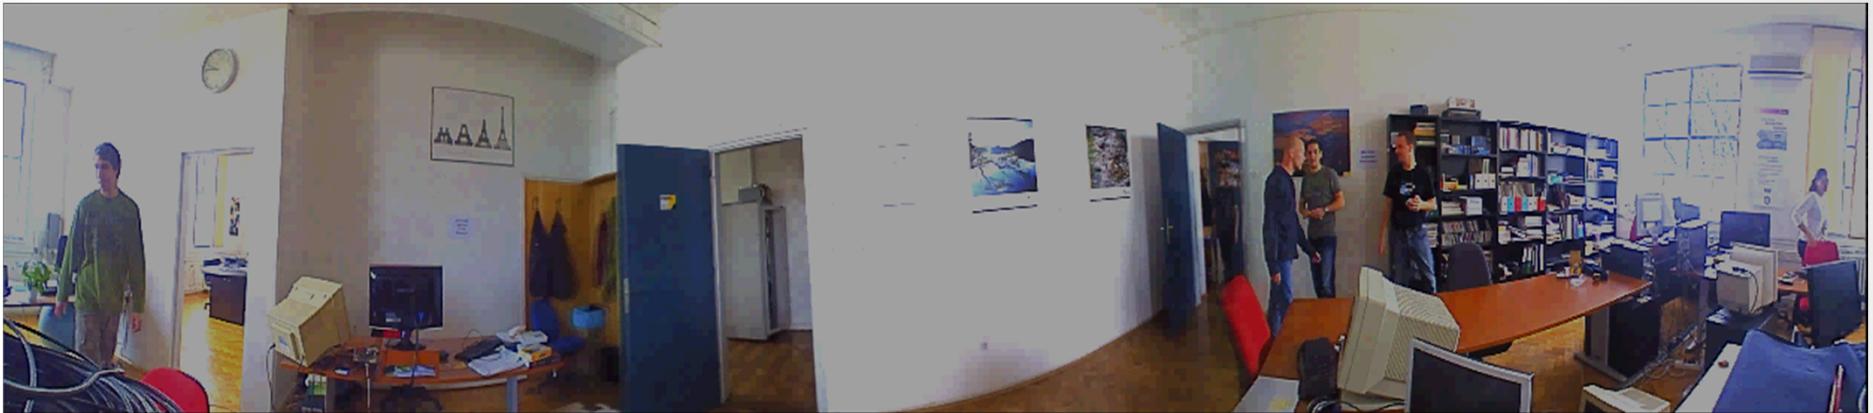
HD Surveillance



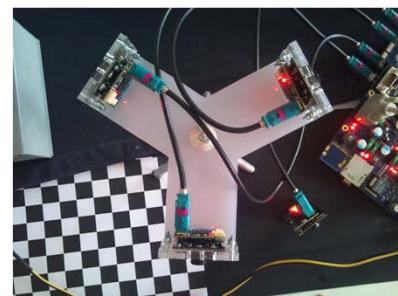
Driver Assist

- Stereo (+ IR) cameras for depth map
- Robots operating in unrestricted environments
- Simultaneous location and mapping
- Immersive environments
- Massive camera scaling
- Multiple camera stitching and tracking
- Semantic understanding and decision making

An Example



3*1 MP Camera Stitching from Xylon
(<http://www.xylon.com>)

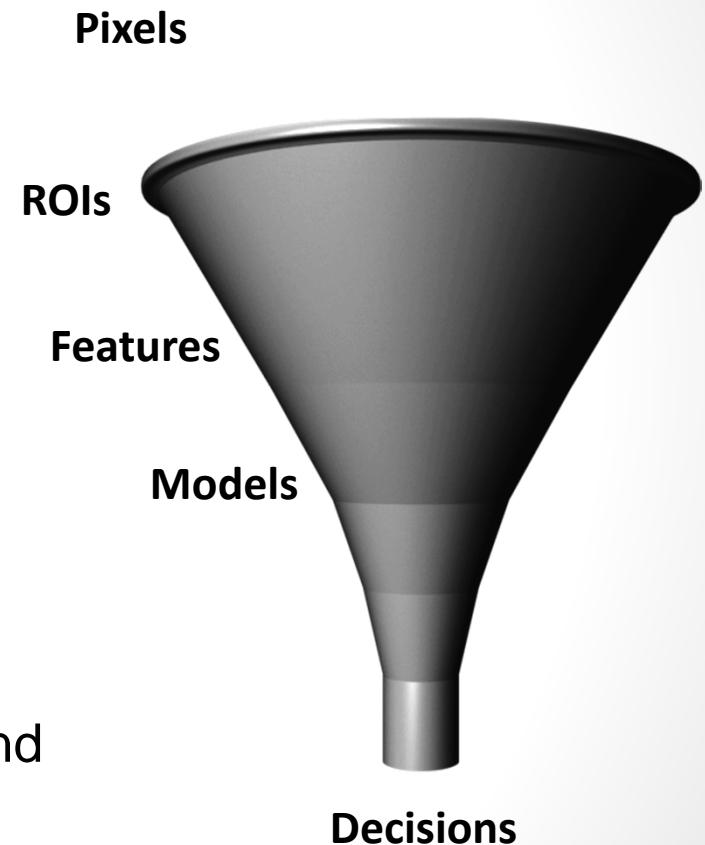


- Collect the right information
 - Fuse 3 cameras into one image
- Detect interesting things
 - There are 4 people
- Identify the things
 - Person in the green shirt is Fred
- Understand the environment
 - Fred is leaving
- Make a decision
 - Turn on the lights in the next room and warm up Fred's car

Similar challenges in many vision applications

The Funnel

- Collect the right information
 - Fuse 3 cameras into one image
- Detect interesting things
 - There are 4 people
- Identify the things
 - Person in the green shirt is Fred
- Understand the environment
 - Fred is leaving
- Make a decision
 - Turn on the lights in the next room and warm up Fred's car



Similar challenges in many vision applications

Video Processing Rates

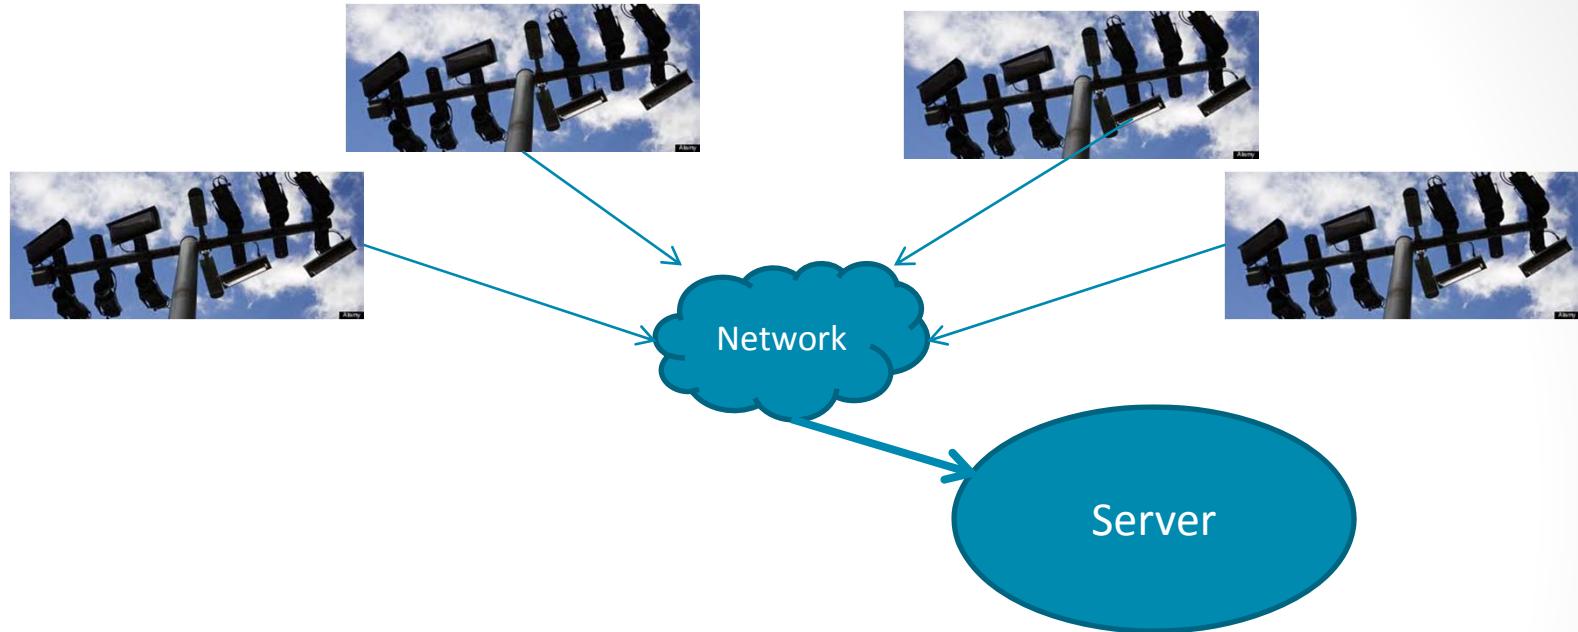
Design approach	RISC Proc.	Proc. w/ accels.	Folded datapath	Pipelined datapath	Replicated datapath
clock:sample	1000:1	100:1	10:1	1:1	1:10
Sample Rate (200MHz clock)	200Ks/s	2Ms/s	20Ms/s	200Ms/s	2 Gs/s
Computation (500 ops/sample)	100M op/s Feature Processing	1G op/s	10G op/s	100G op/s FullHD Pixel rate Processing	1 T op/s



Arm A9 processors
1-2 Gops each

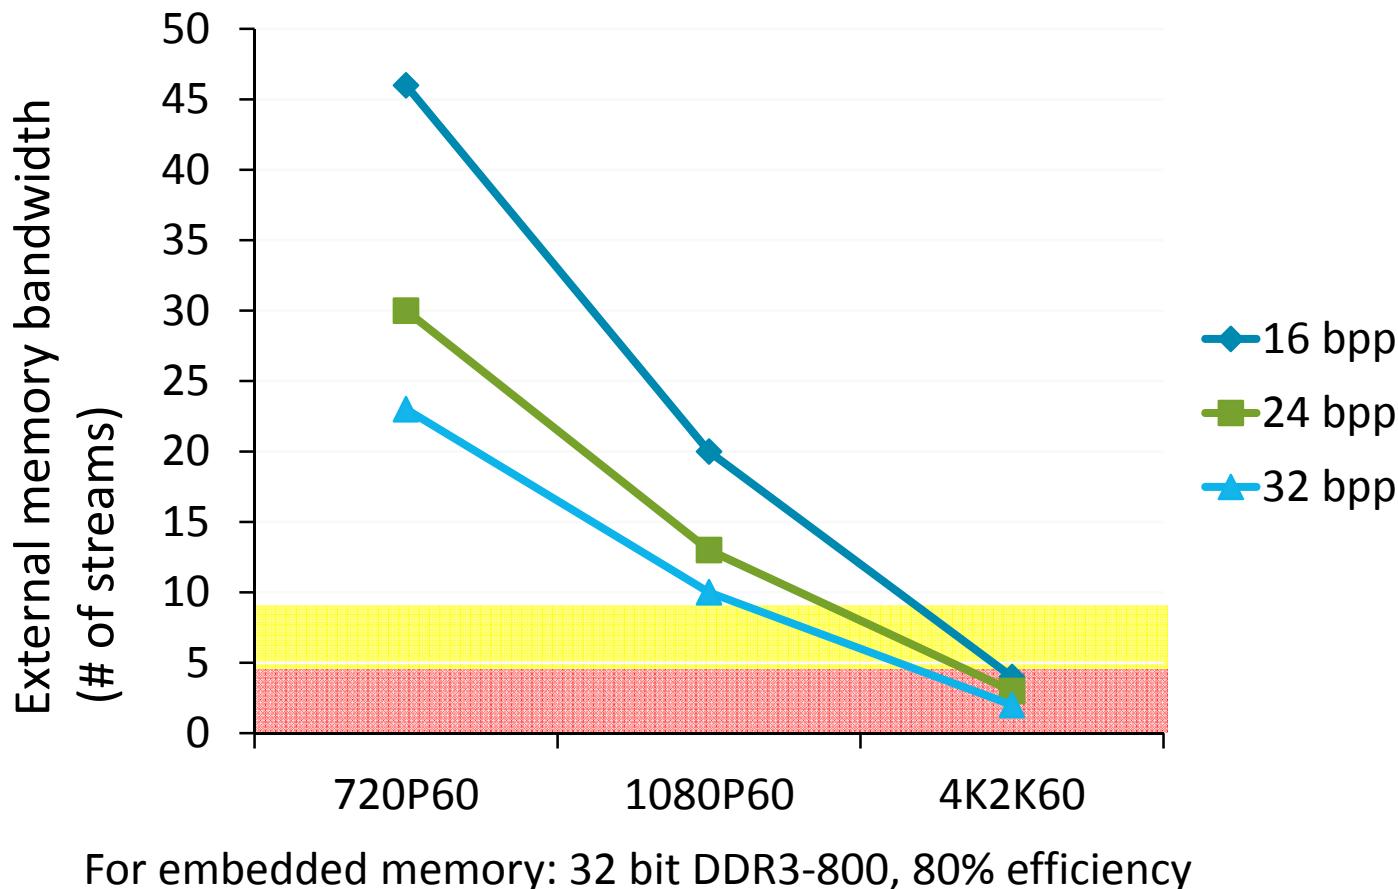
specialized hardware
10 – 500 Gops

In the network?



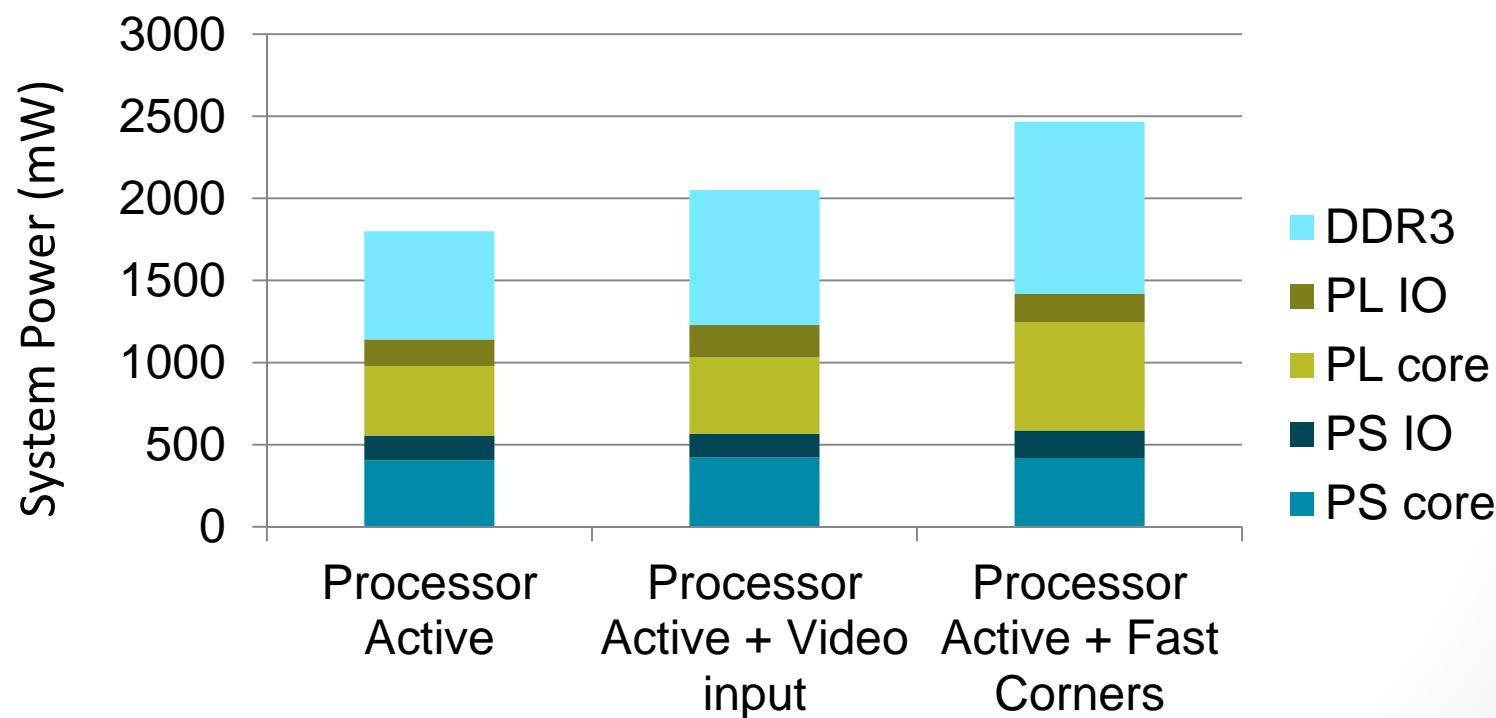
- 1080P60 \approx 125 Mpixels/sec * 8 bpp \approx 1Gbit/sec
 - Completely saturates a Gigabit Ethernet link
- Some solutions
 - 'Exotic' networks (10GigE)
 - Point-to-point (SDI/HDMI)
 - Compression (H264/H265)

Video Streams to Memory



Process and decide without storing whenever possible!

- External Memory Power can also be significant
- Real-time HD video processing can be very efficient in FPGAs (200-300 mW incremental power)



Data collected from Xilinx Application Note 1157

© Copyright 2014 Xilinx

Analytics Challenges

Network Challenge

- More data than can be stored or sent
- Compression vs. Storage
- Intelligent Recording

Power Challenge

- Power over Ethernet
- Thermal Issue
- Power Management

Cost Challenge

- Existing infrastructure
- Increase in camera types
- New applications without new hardware

Dumb Cameras ->
Smarter Cameras

Efficient
Embedded
Processing

Increasing
Integration with
programmability

OpenCV Overview

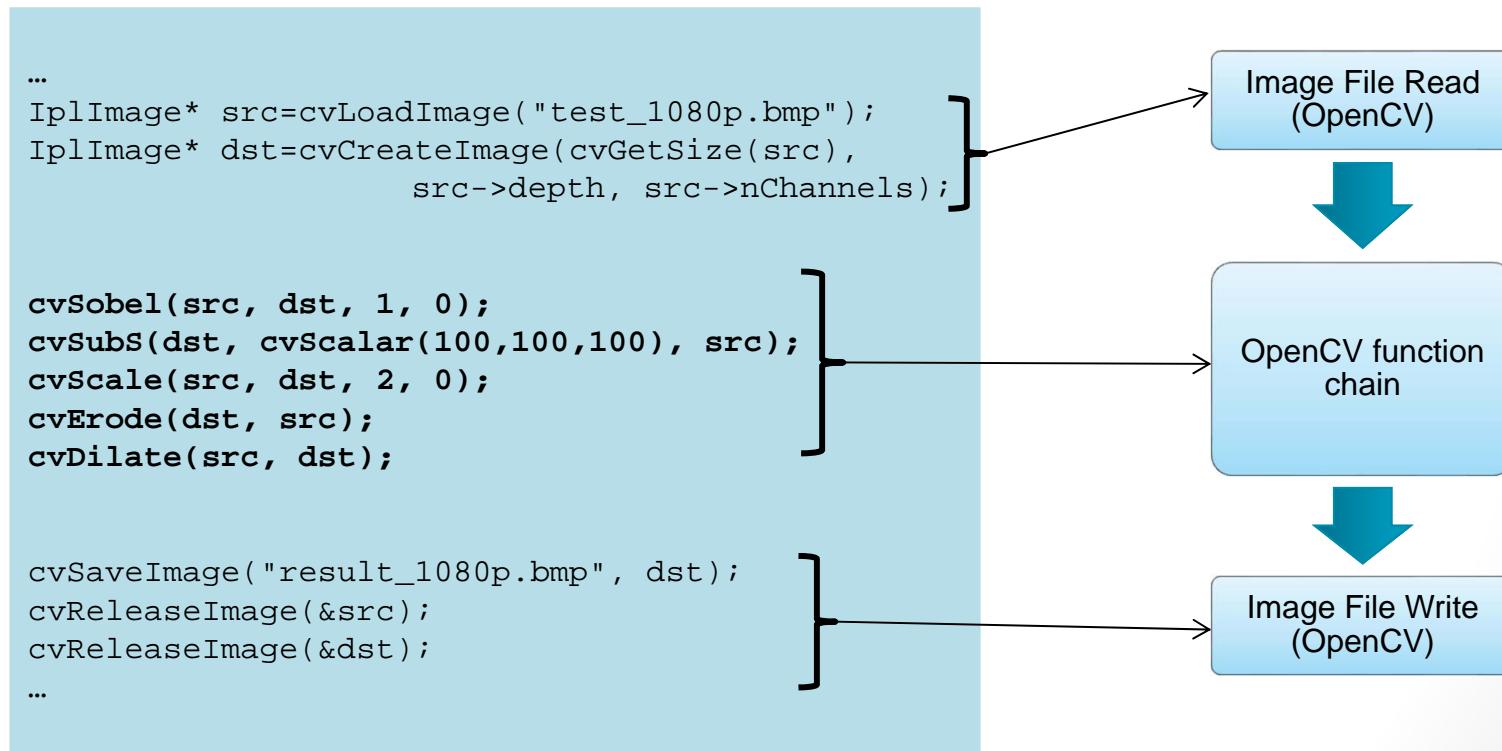
- Open Source Computer Vision (OpenCV) is widely used to develop computer vision applications
 - Library of 2500+ optimized video functions
 - Optimized for desktop processors and GPUs
 - Tens of thousands of users
- Big Questions
 - How to get these applications into an embedded system?
 - What is right mix of processing for different parts of an application?



OpenCV modules

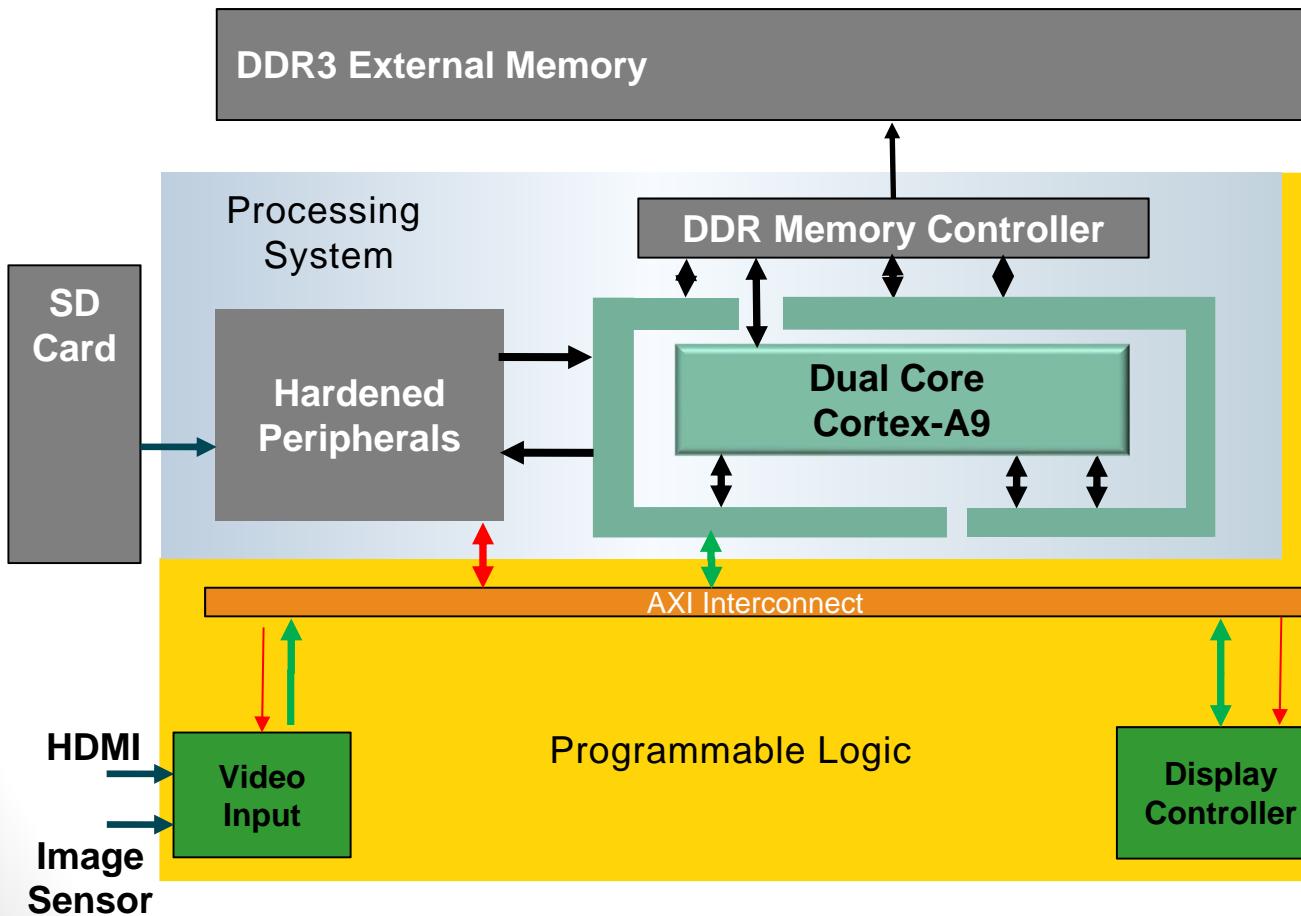
- core - key data structures
- imgproc - image filtering
- features2d - corner, blob extraction and matching
- stitching - lens correction, image stitching, stereo rectification
- objdetect - classifiers
- video - optical flow
- videostab - video stabilization
- flann - flann feature matching
- ml - machine learning, classifier training
- calib3d - lens calibration and stereo alignment
- highgui - platform independent GUI toolkit

- One image input, one image output
 - Processed by chain of functions sequentially



Example Architecture

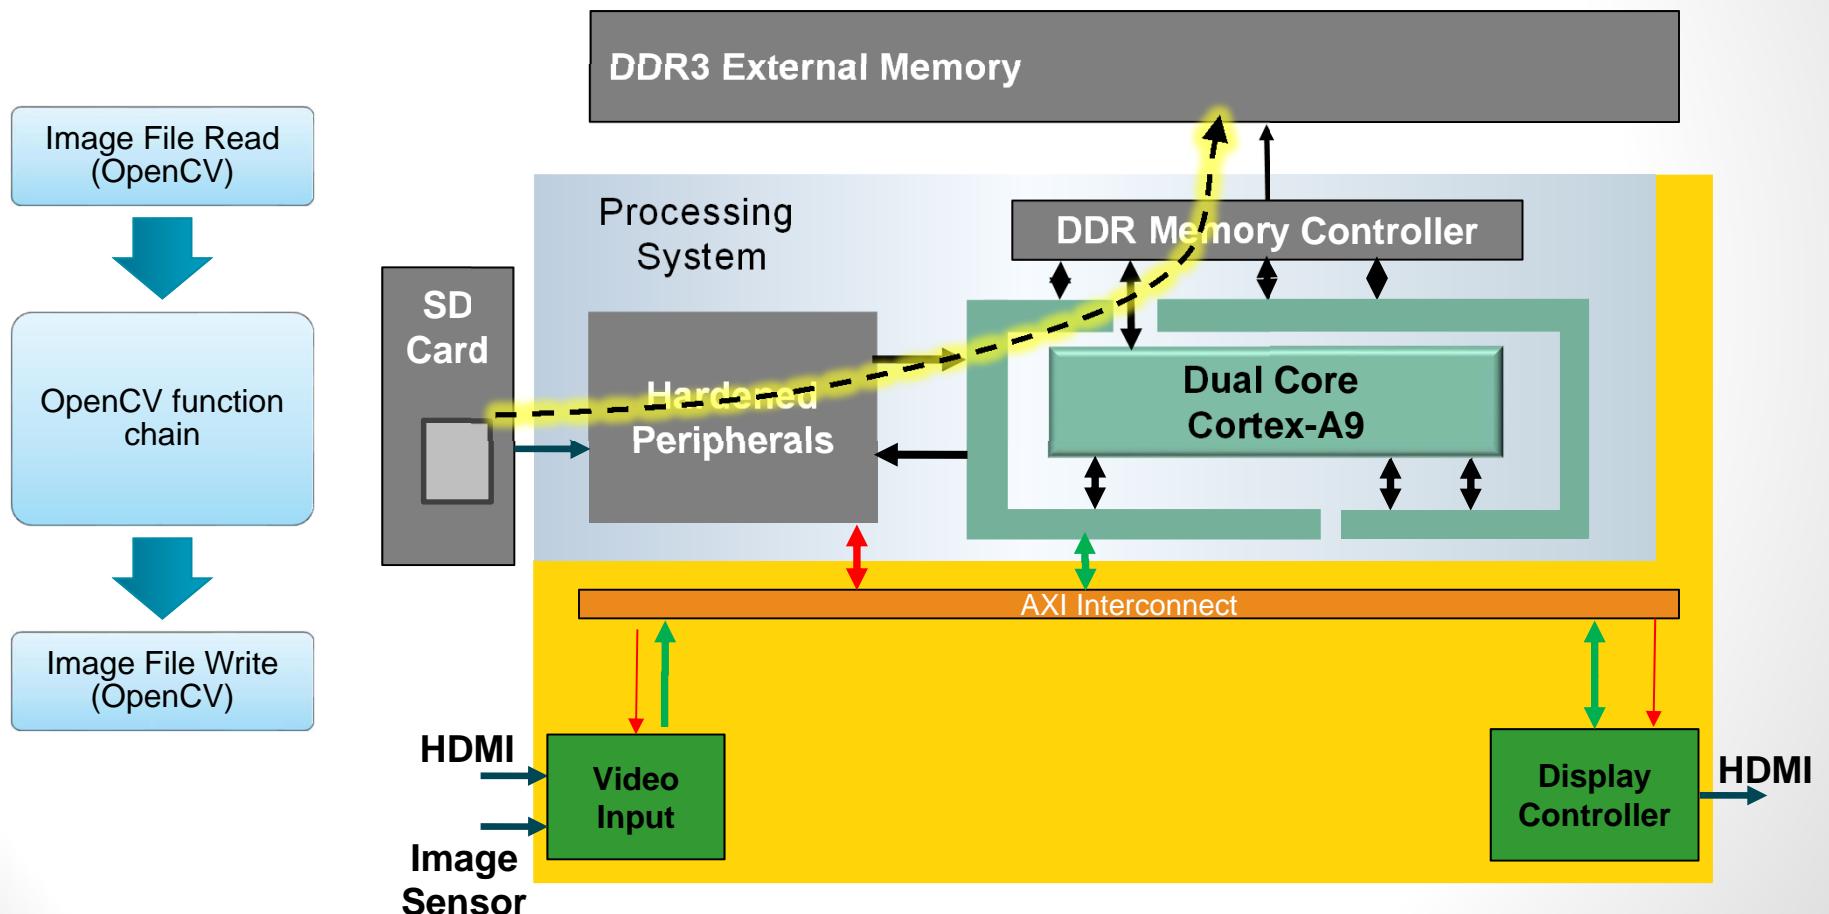
- Xilinx 'Zynq' Video and Imaging Kit
 - Supports HDMI I/O and VITA2000 Image Sensor

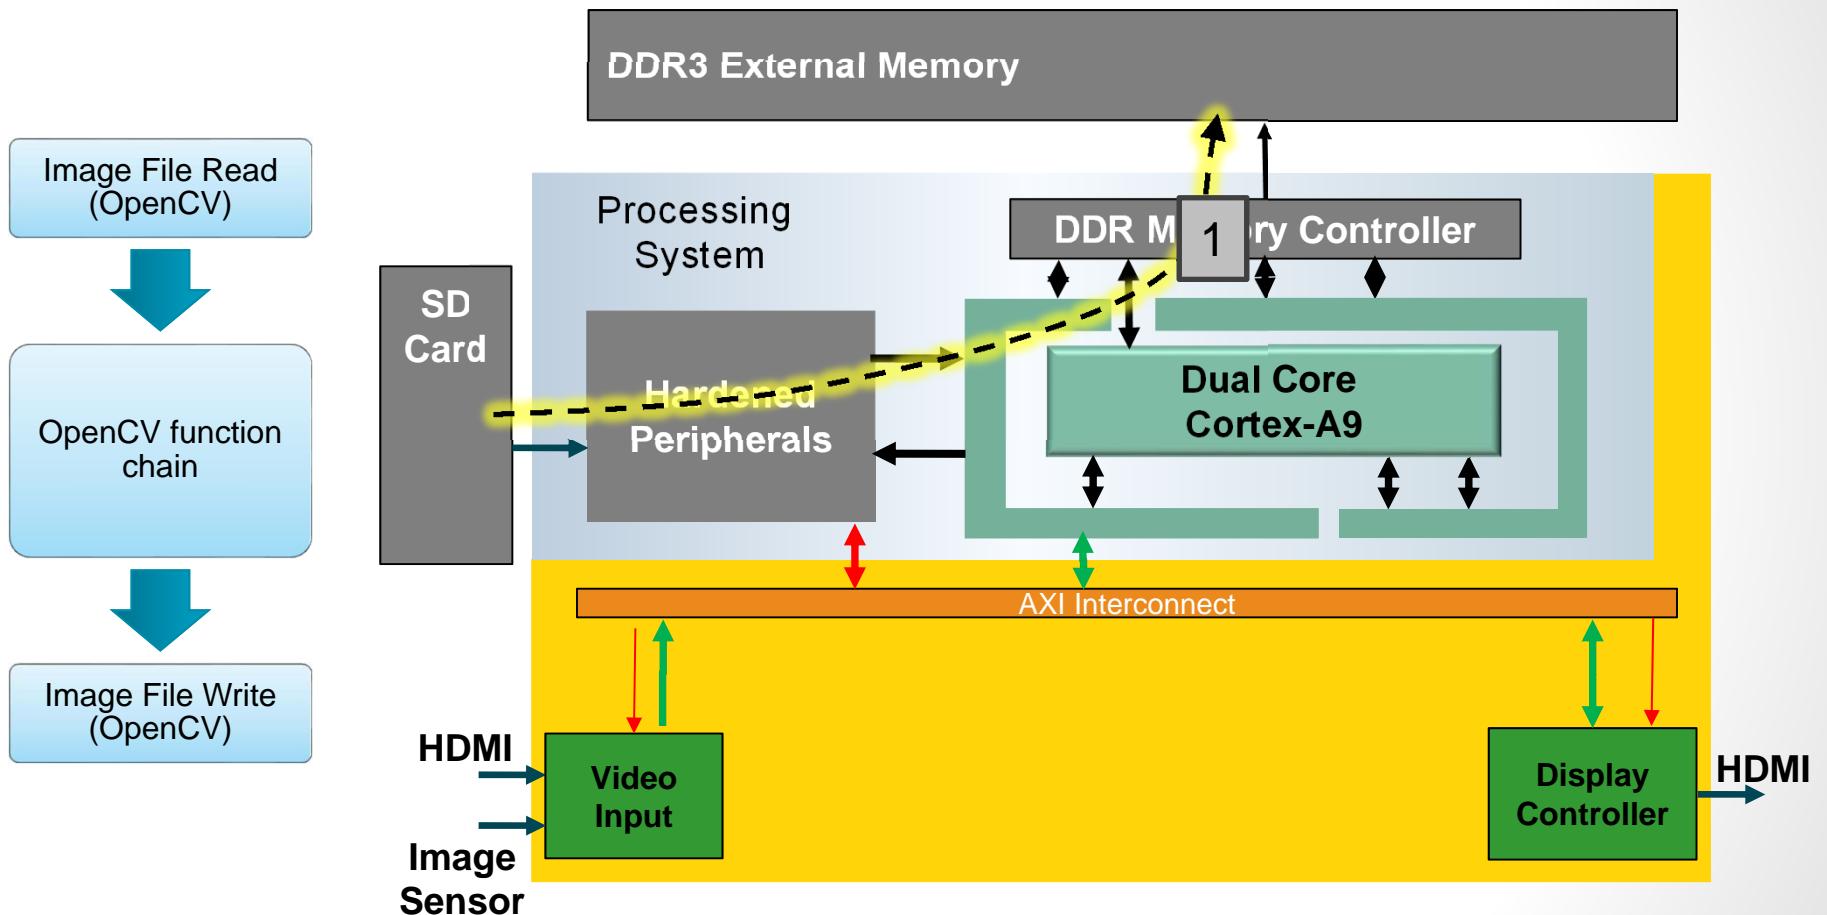


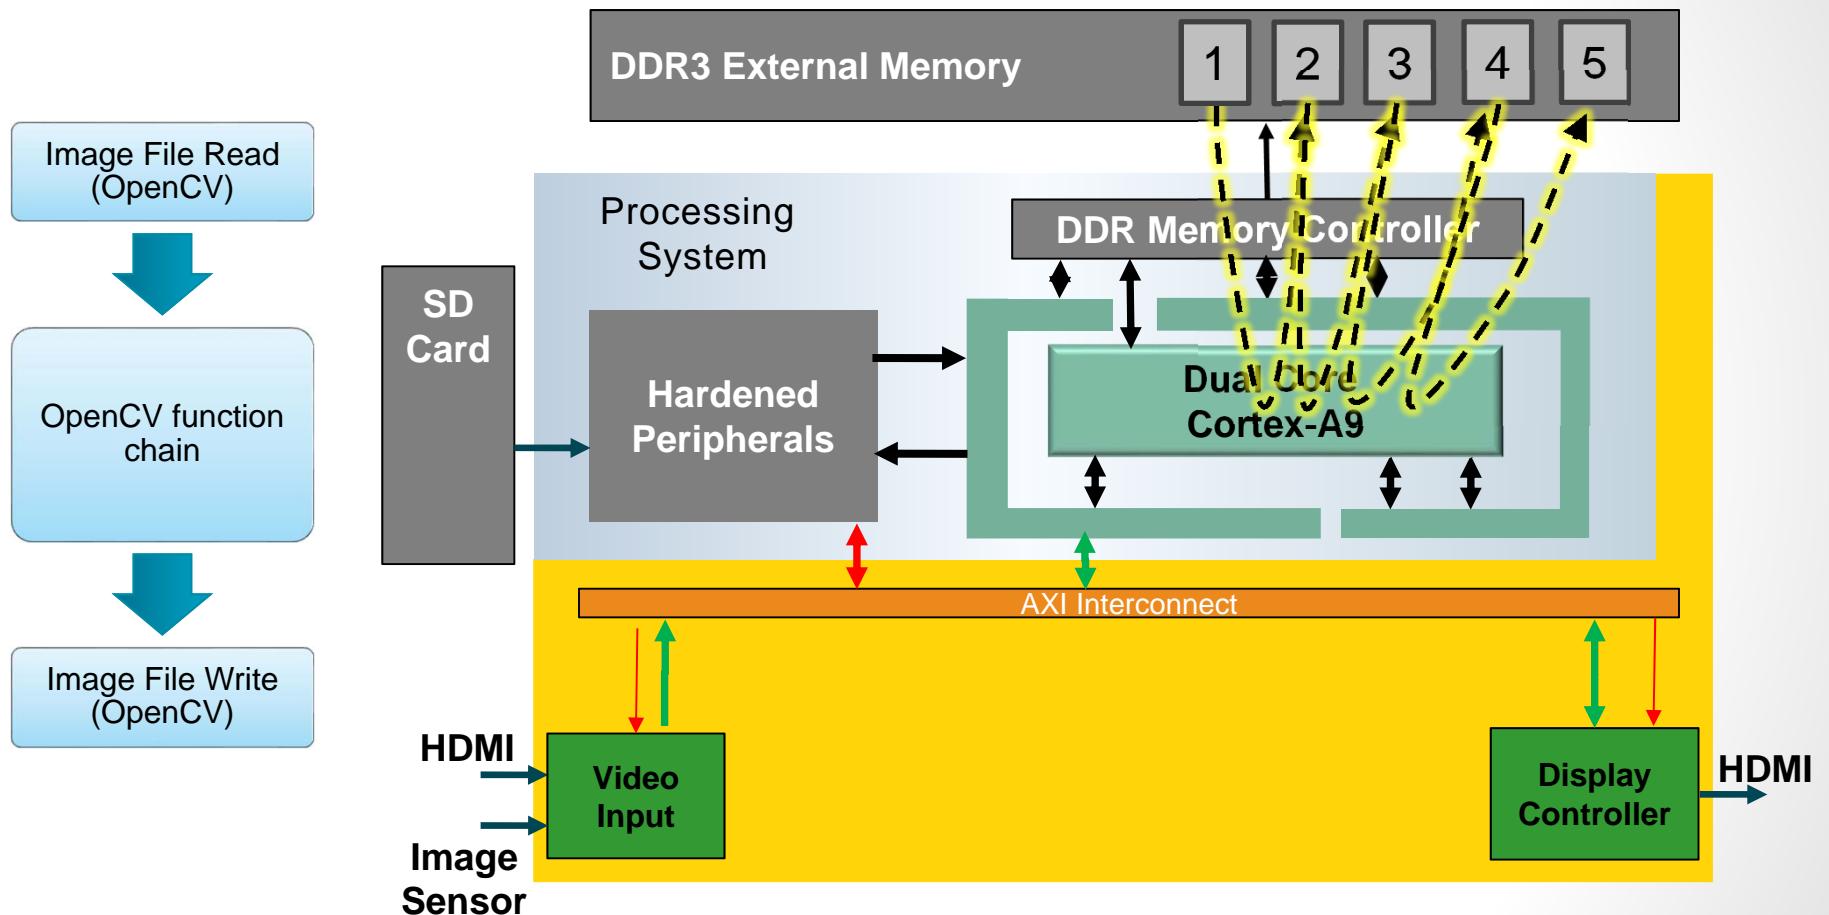
<http://www.xilinx.com/zynq>

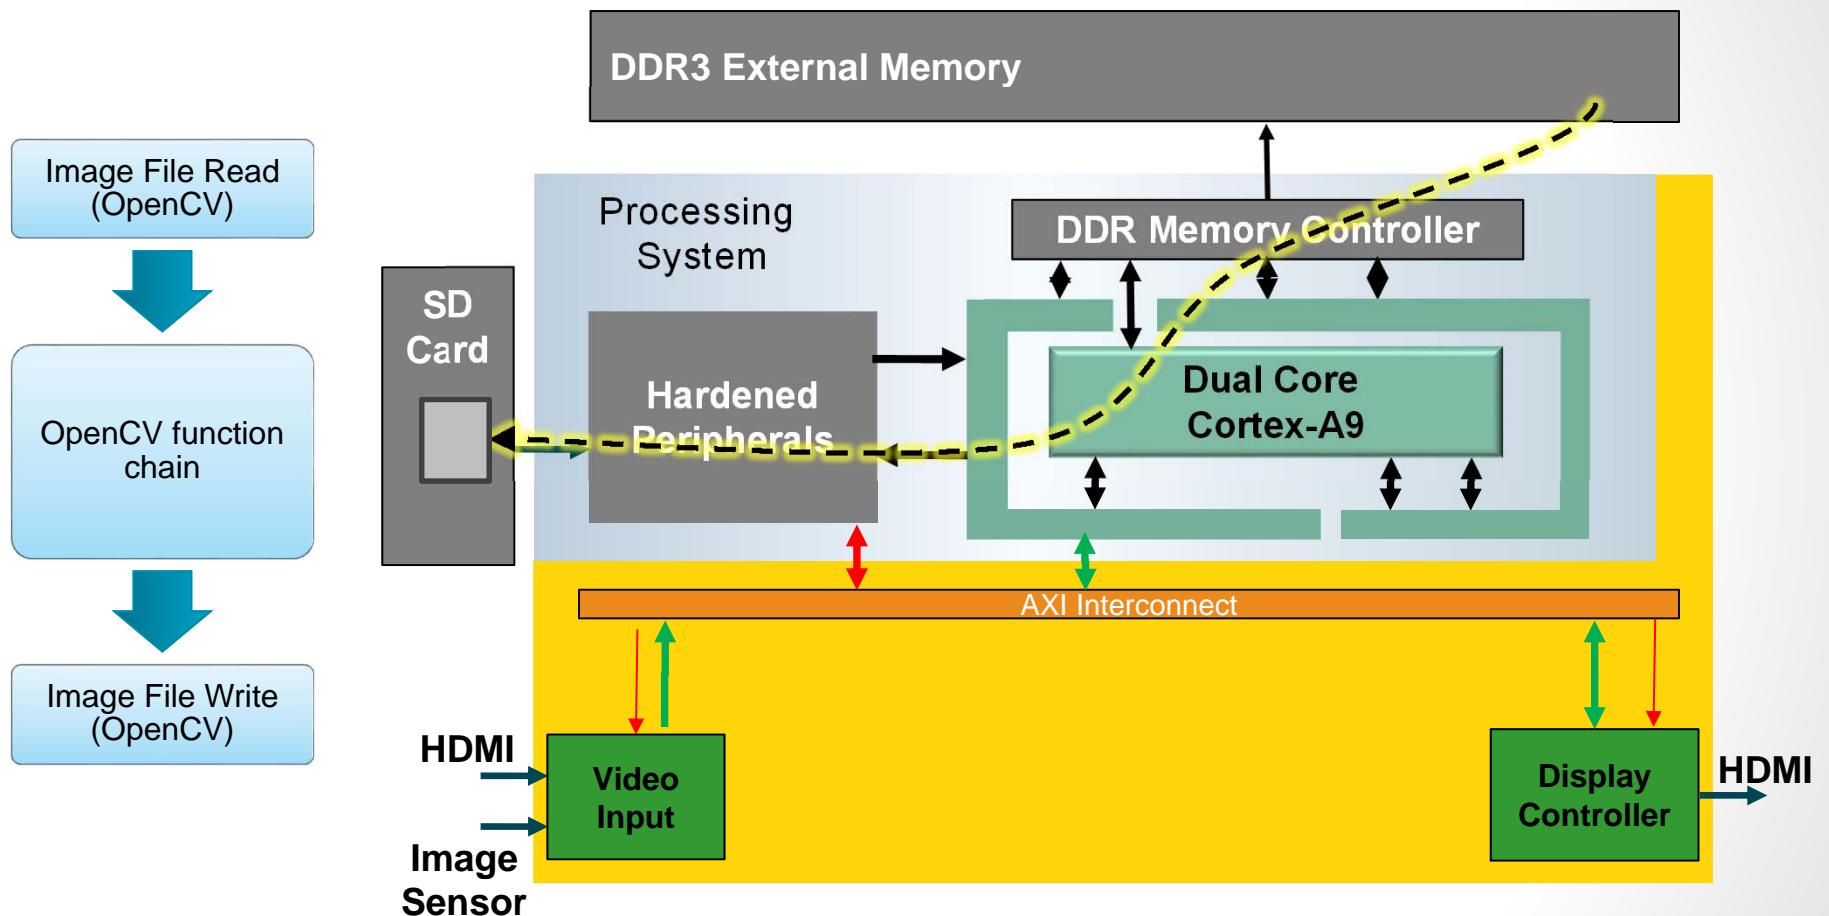
OpenCV on ARM

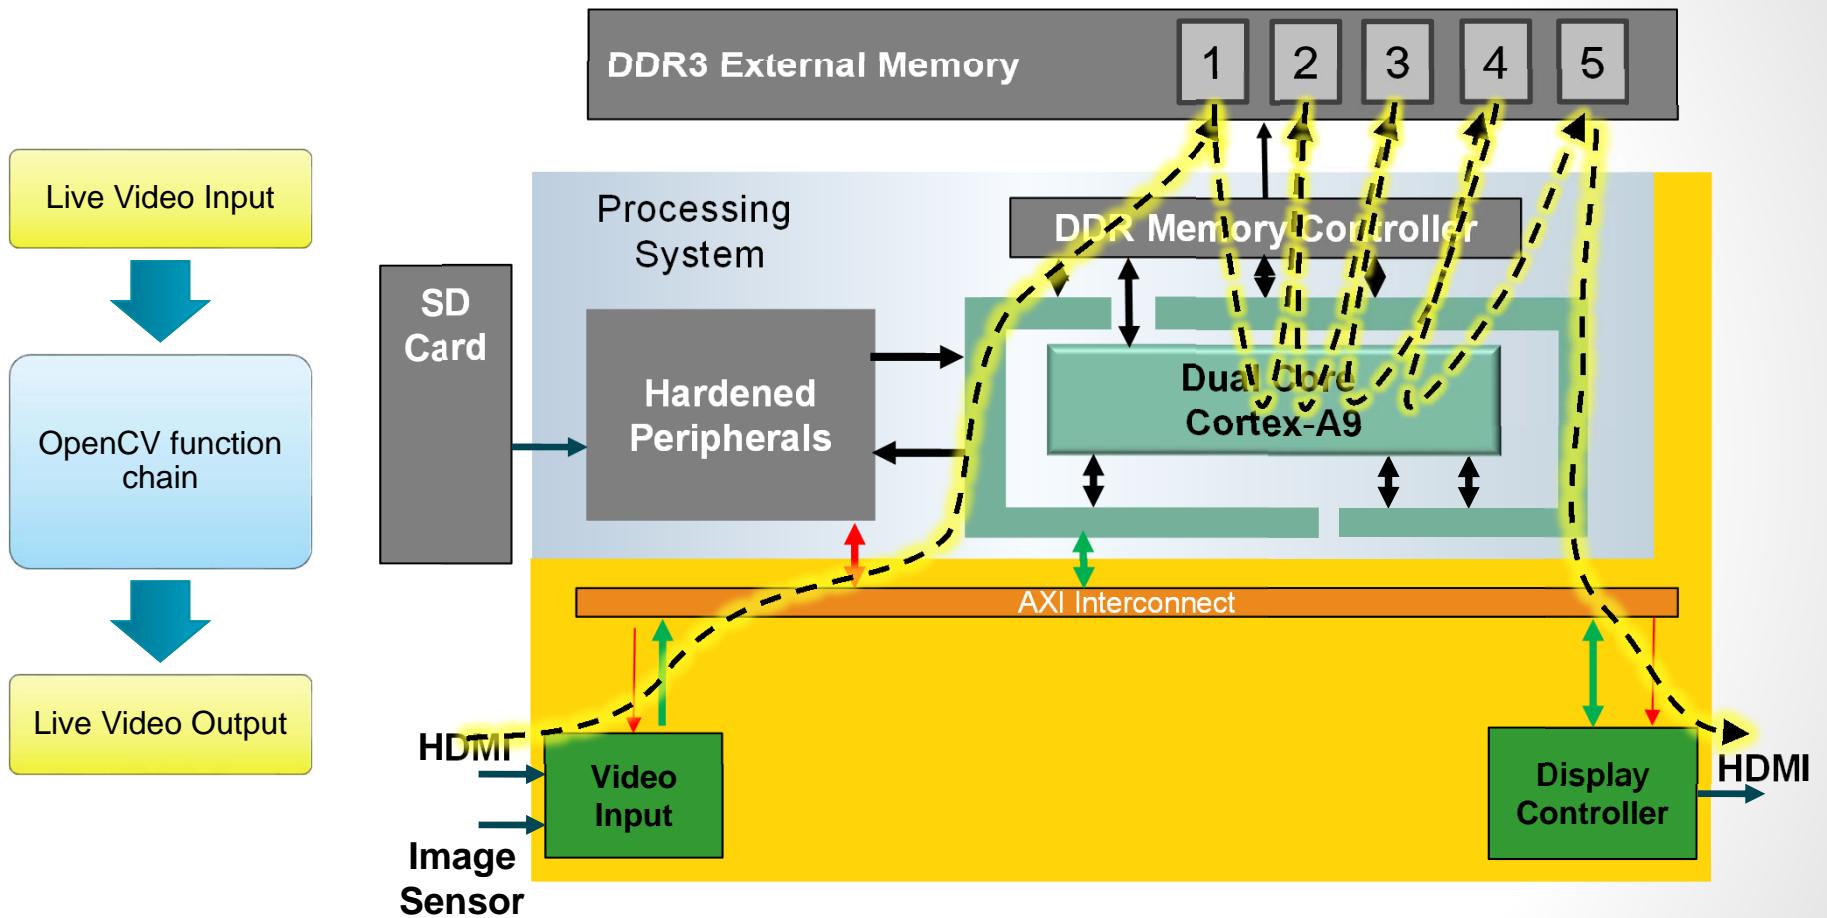
- OpenCV-based image processing built around sequential processing of frame buffers in external memory











ARM Optimized CV

- Explicit prefetching and use of ARM NEON instructions can give significant speedup over generic code
- Throughput drops with composition of functions

Performance of UncannyCV library (<http://www.uncannyyvision.com>)

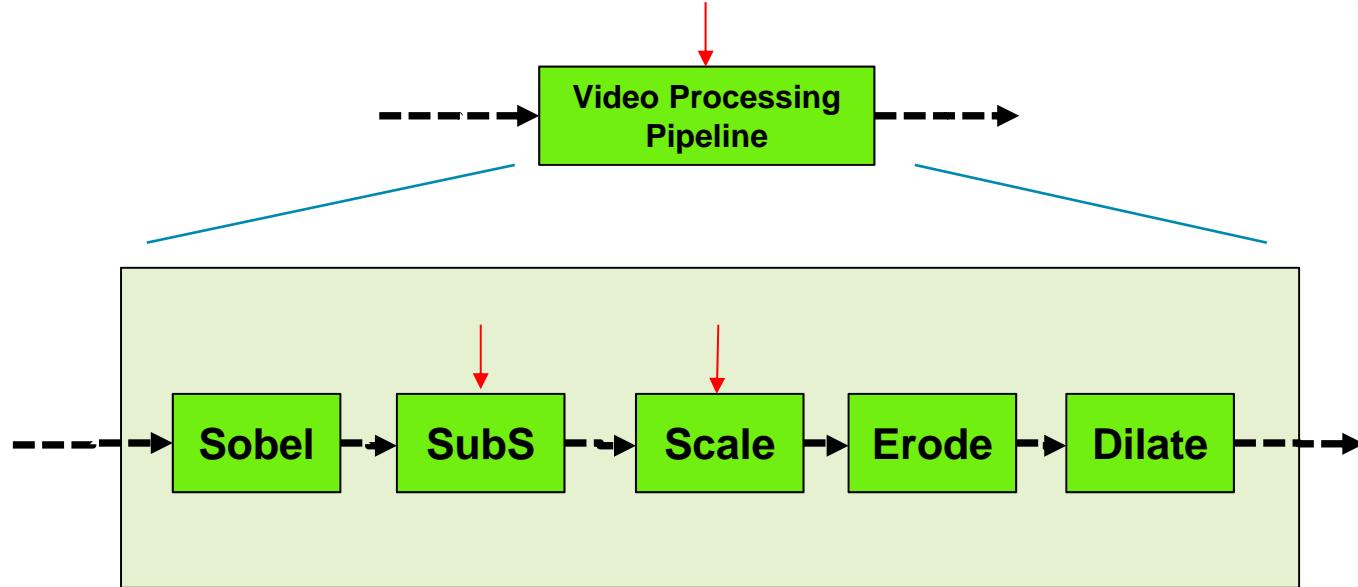
	Throughput * (Megapixels per sec)	Acceleration * (vs. OpenCV on ARM)
FAST9 Corner Detection	24	2x
Canny Edge Detection	25	3x
Harris Corner Detection	15.7	6.5x
Erosion/Dilation	153	6.5x
5x5 convolution	96	22x
LBP Face Detection		3x
HOG Pedestrian Detection		9x

* based on single core 1GHz Cortex-A9 (2 wide SIMD)

ARM Limitations

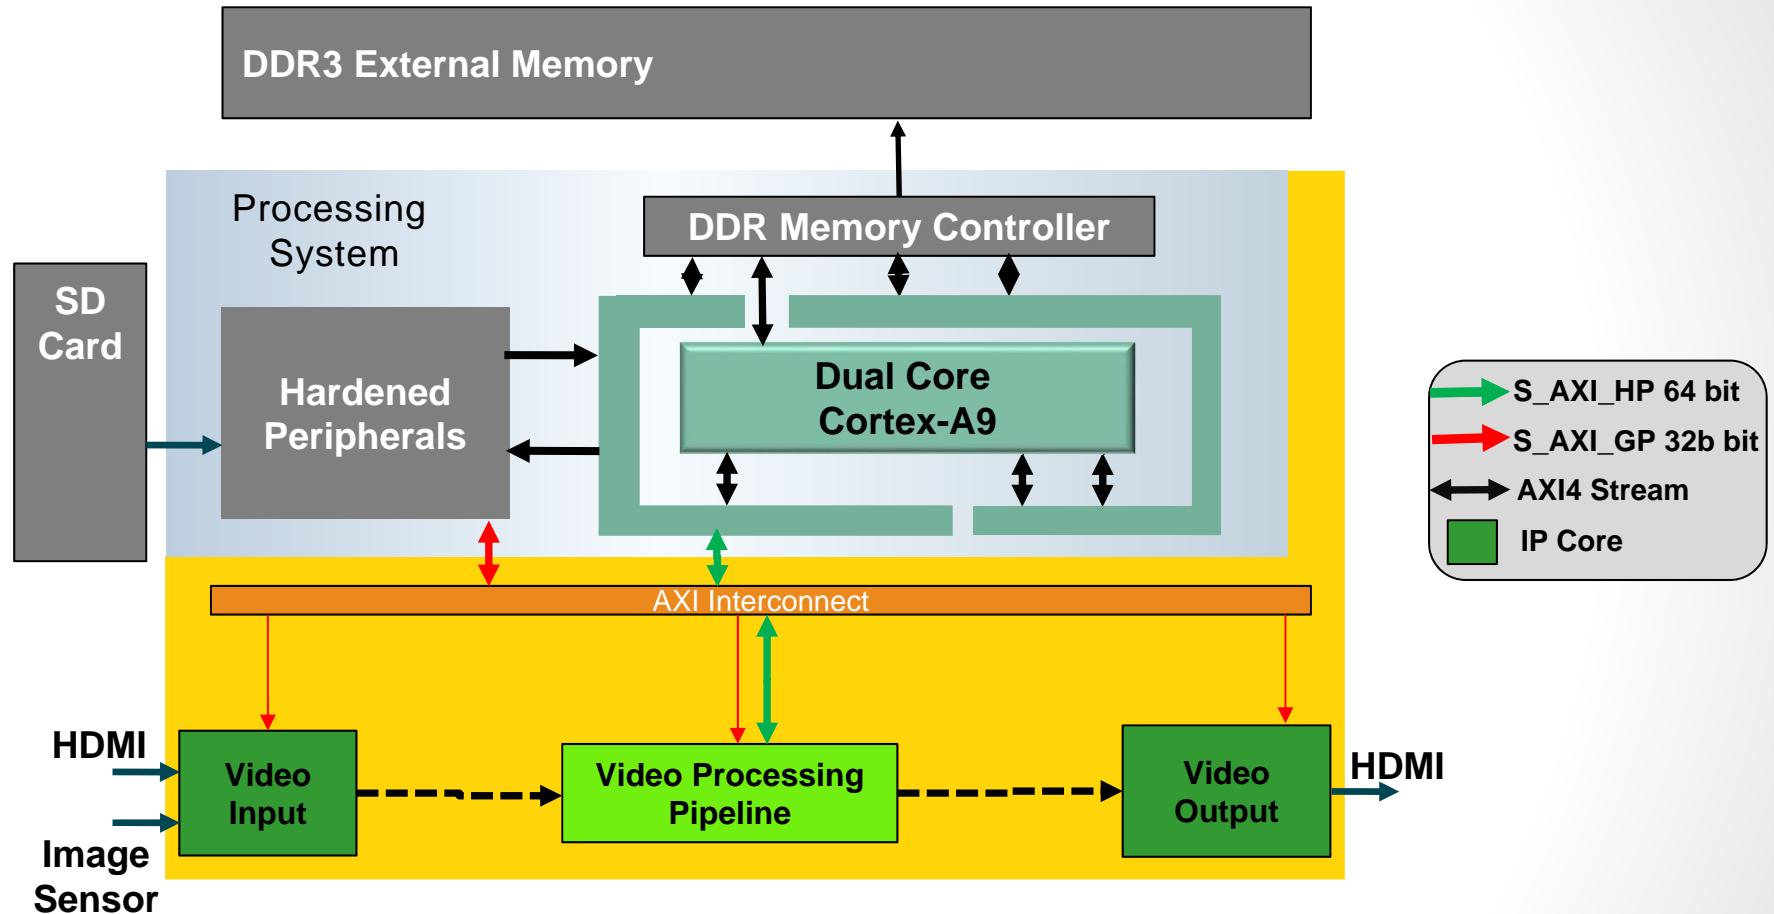
- Poor access locality -> small caches perform poorly
 - 720P \approx 1 M pixels/frame \approx 4 MByte/frame @ 32 bpp
 - Need multiple frames in cache for best performance
 - Need good cache prefetching performance
- Generic processors limited in parallel operations/number of cores.
 - 720P60 \approx 60 Mpixels/sec \approx 16 cycles per pixel @ 1 Ghz
 - Complex algorithms need many operations per cycle.

Real-time HD processing requires solving both problems



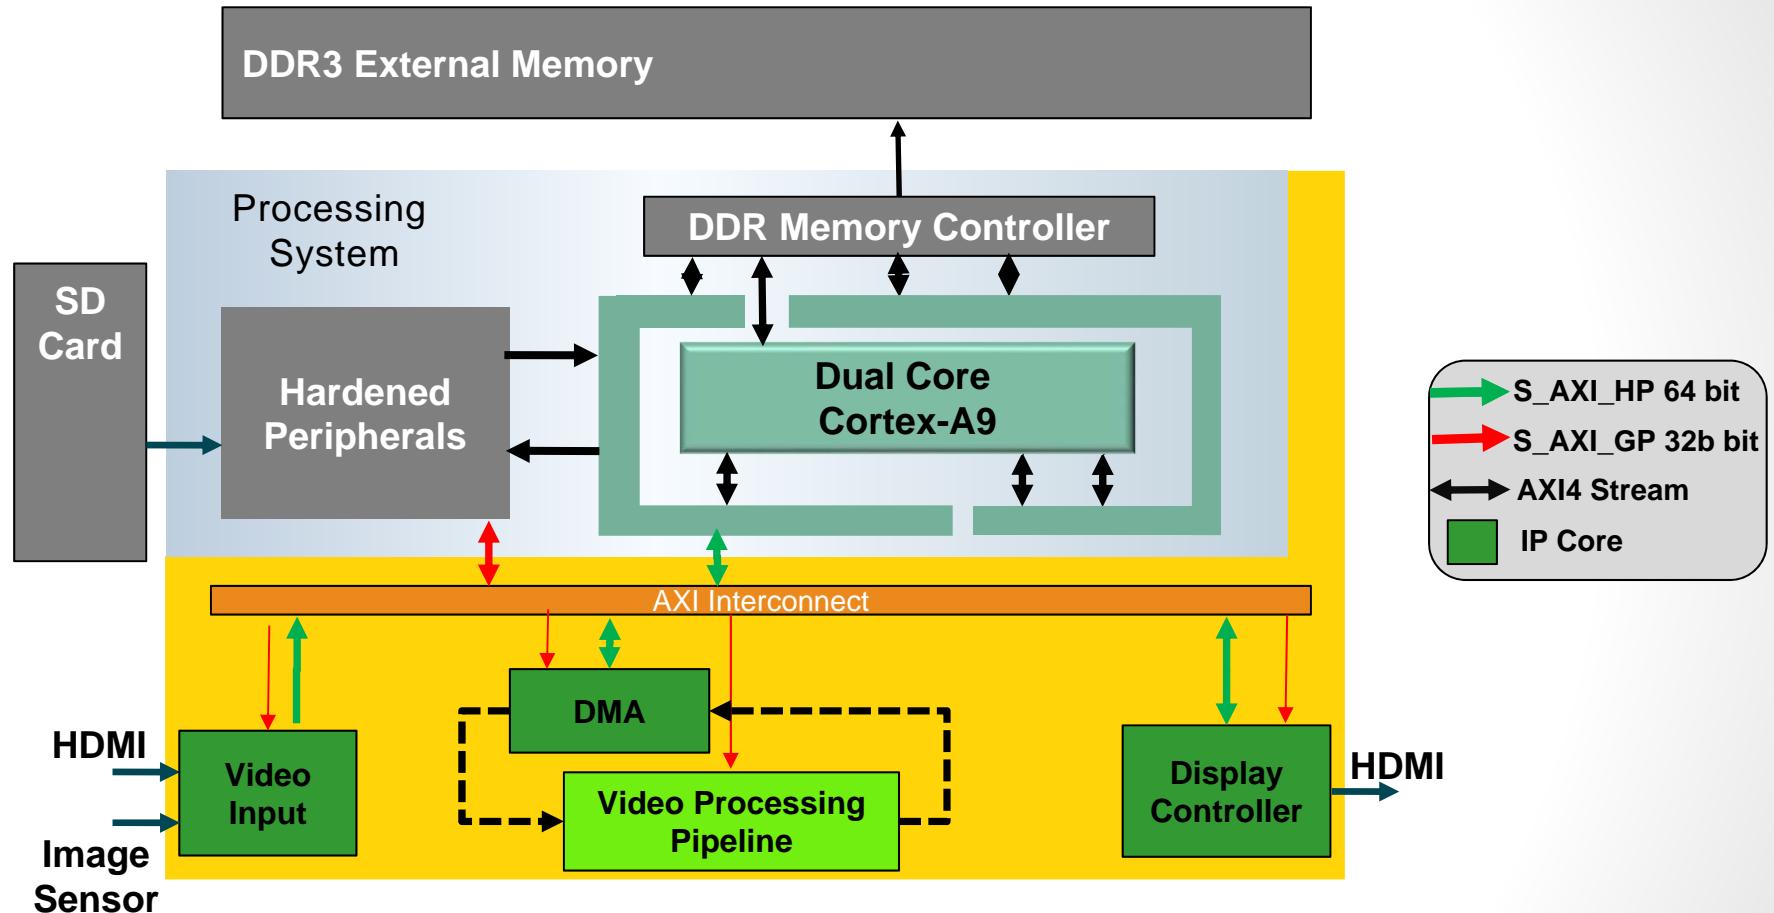
- Avoids need for intermediate buffers
 - Pixels (or tiles of pixels) can be passed directly from one block to another
- Scales well to complex algorithms
 - Additional processing need not share computing resources

Embedded Vision HW



Pro: Reduced memory bandwidth

Con: Less flexibility and ability to interact with the processor

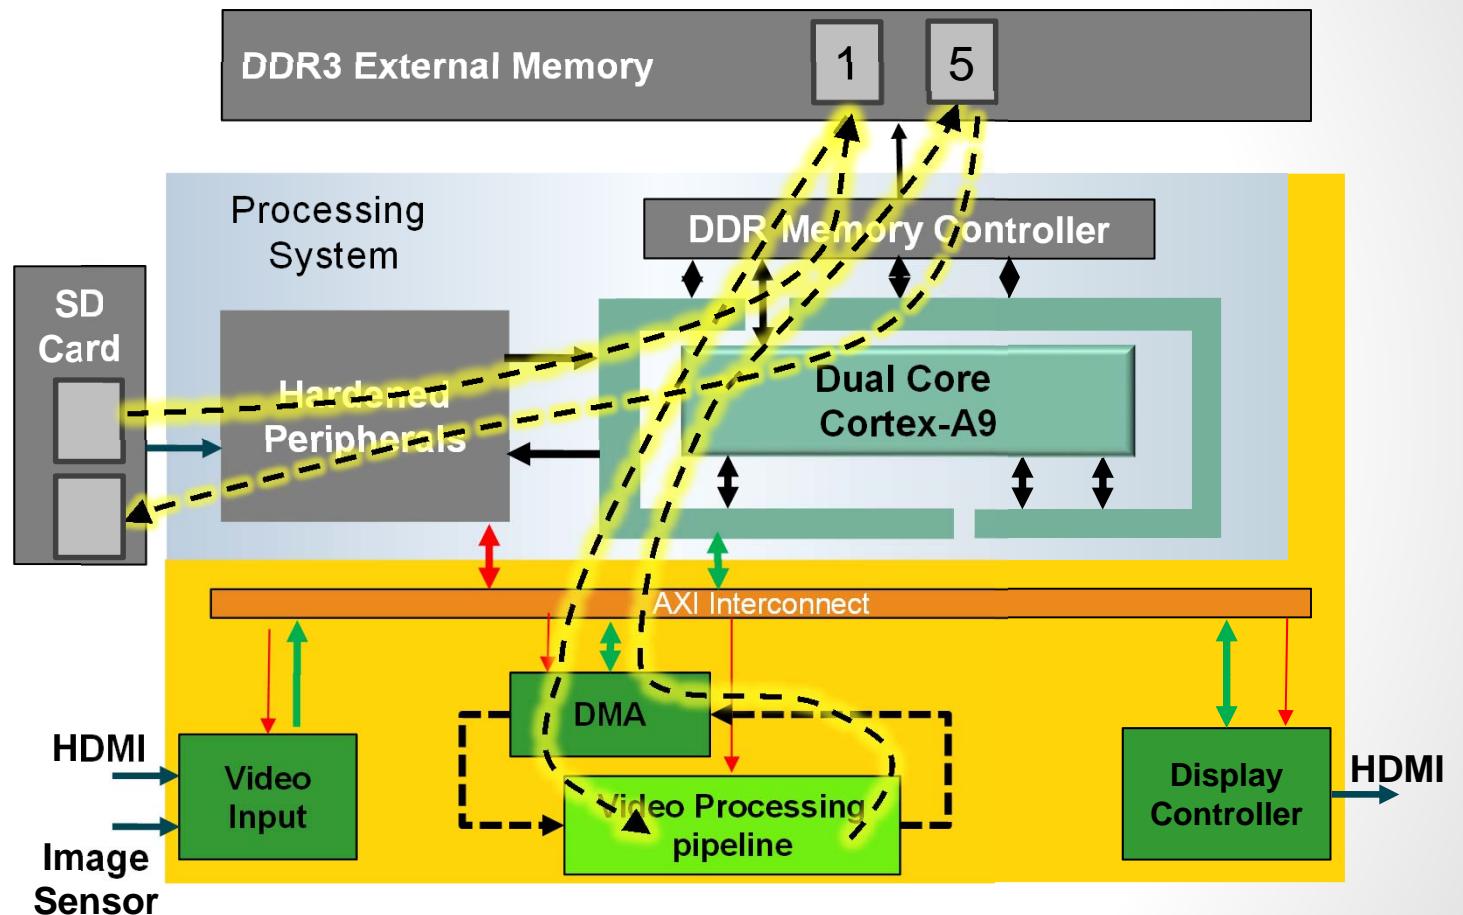
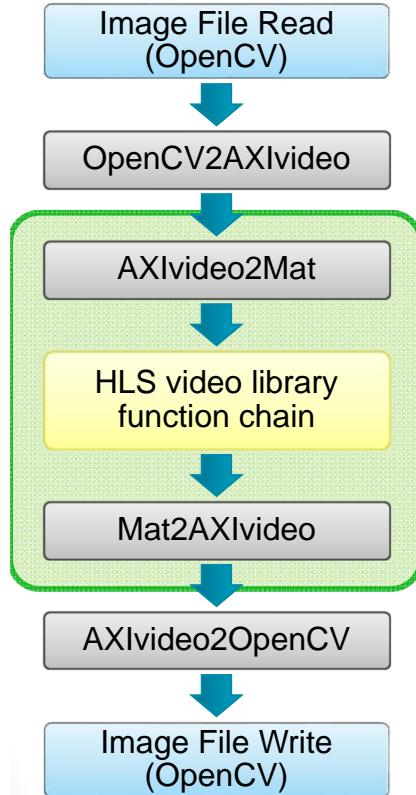


Pro: Processor manages buffers

Con: Requires 4 video streams to external memory

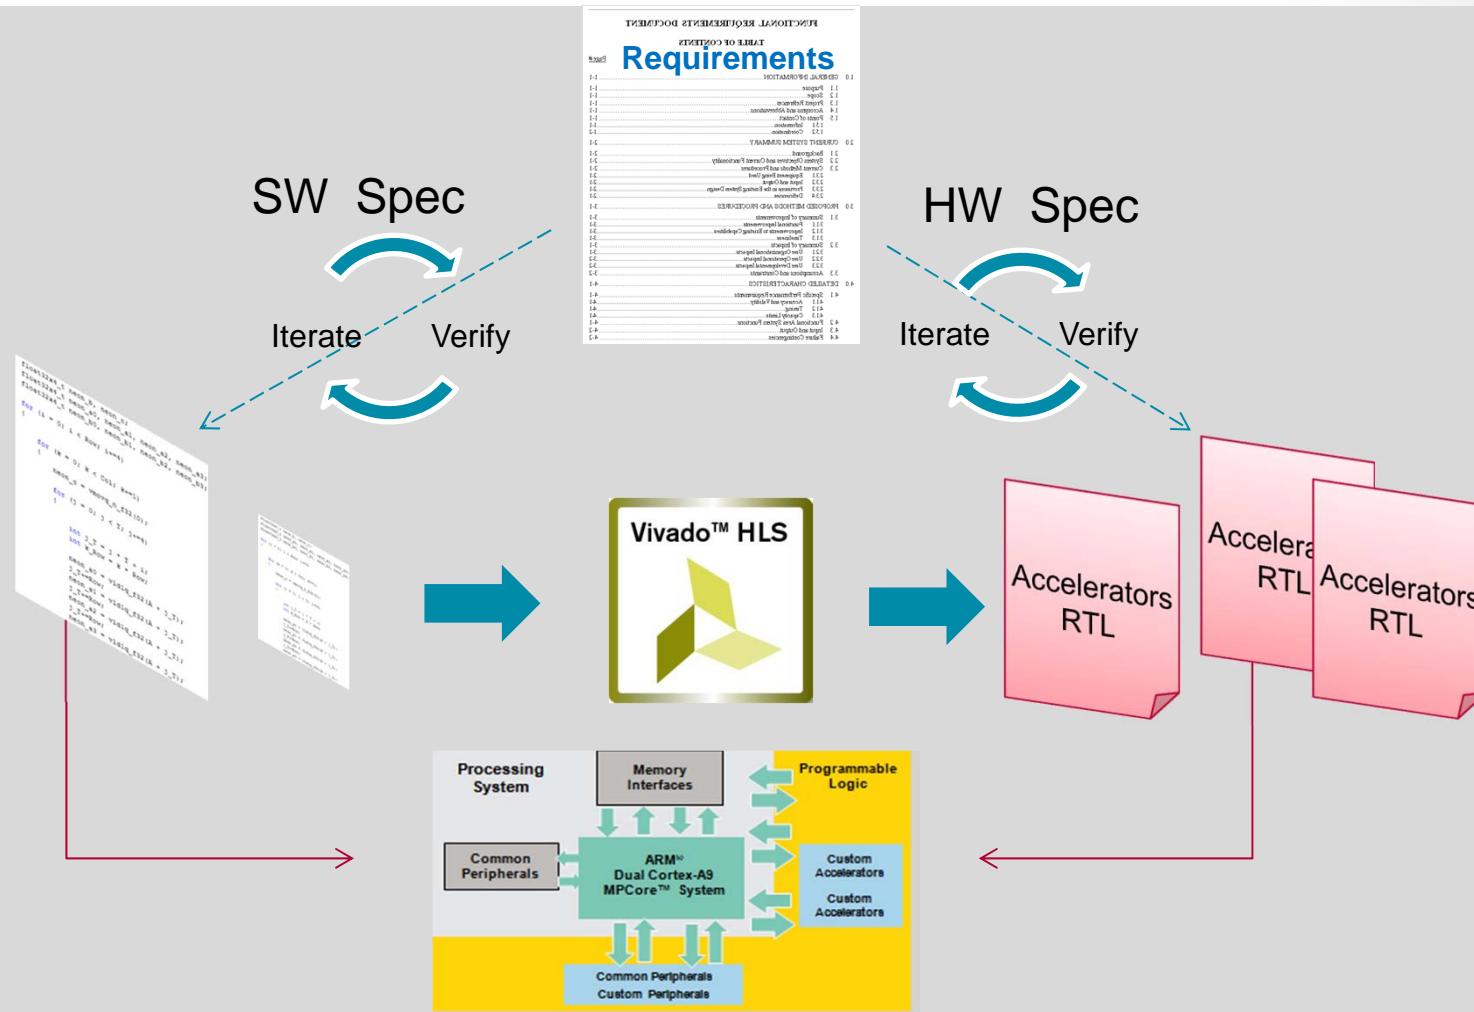
e.g., Zynq Video TRD: <http://www.wiki.xilinx.com/Zynq+Base+TRD+2013.4>

Accelerating Video Processing



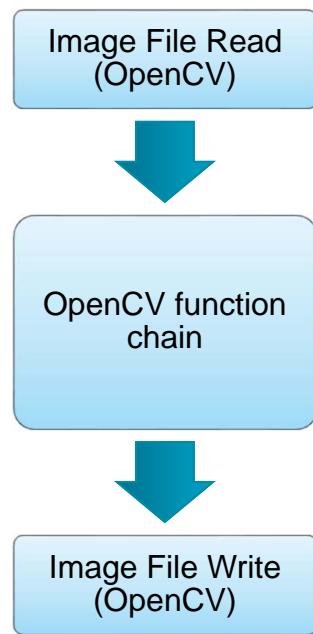
- Most pixel processing algorithms can be chained together with minimal buffering
 - e.g., most filters
 - Each buffer written and read by exactly one function
- Some algorithms require some buffering, but much less than a frame
 - Lens distortion correction
 - Stereo rectification
 - Processing with reconvergent paths
- Some algorithms really require frame buffers
 - Algorithms that process multiple frames (e.g., optical flow)
 - Algorithms with state (e.g., background subtraction)

Using High-Level Synthesis

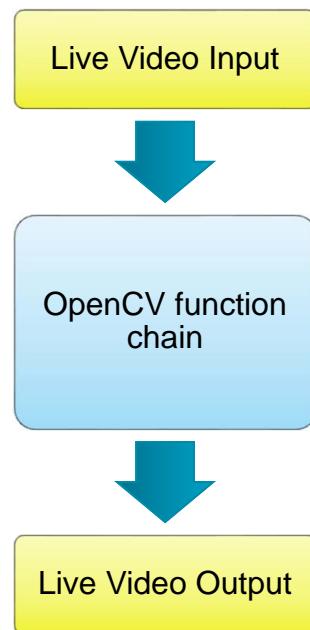


Accelerates Algorithmic C to Co-Processing Accelerator Integration

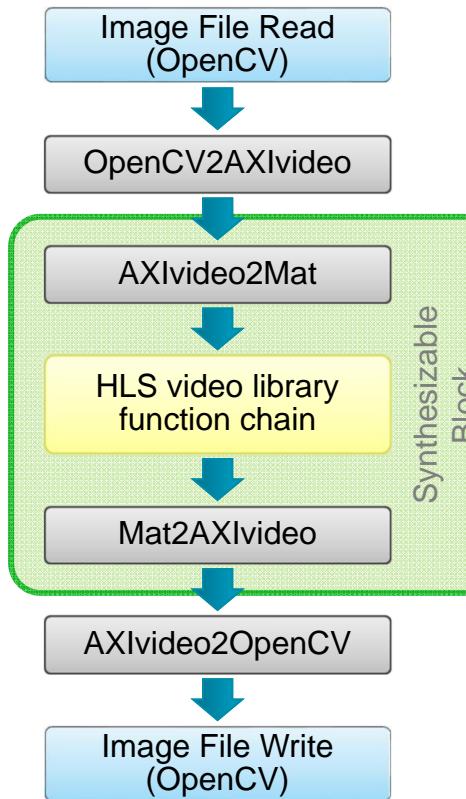
Pure OpenCV Application



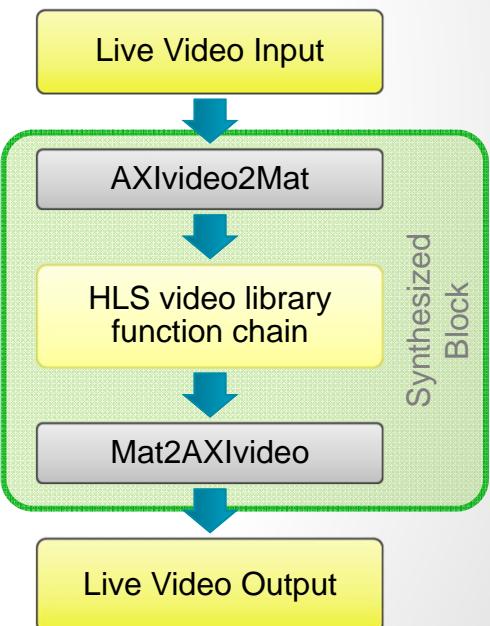
Integrated OpenCV Application

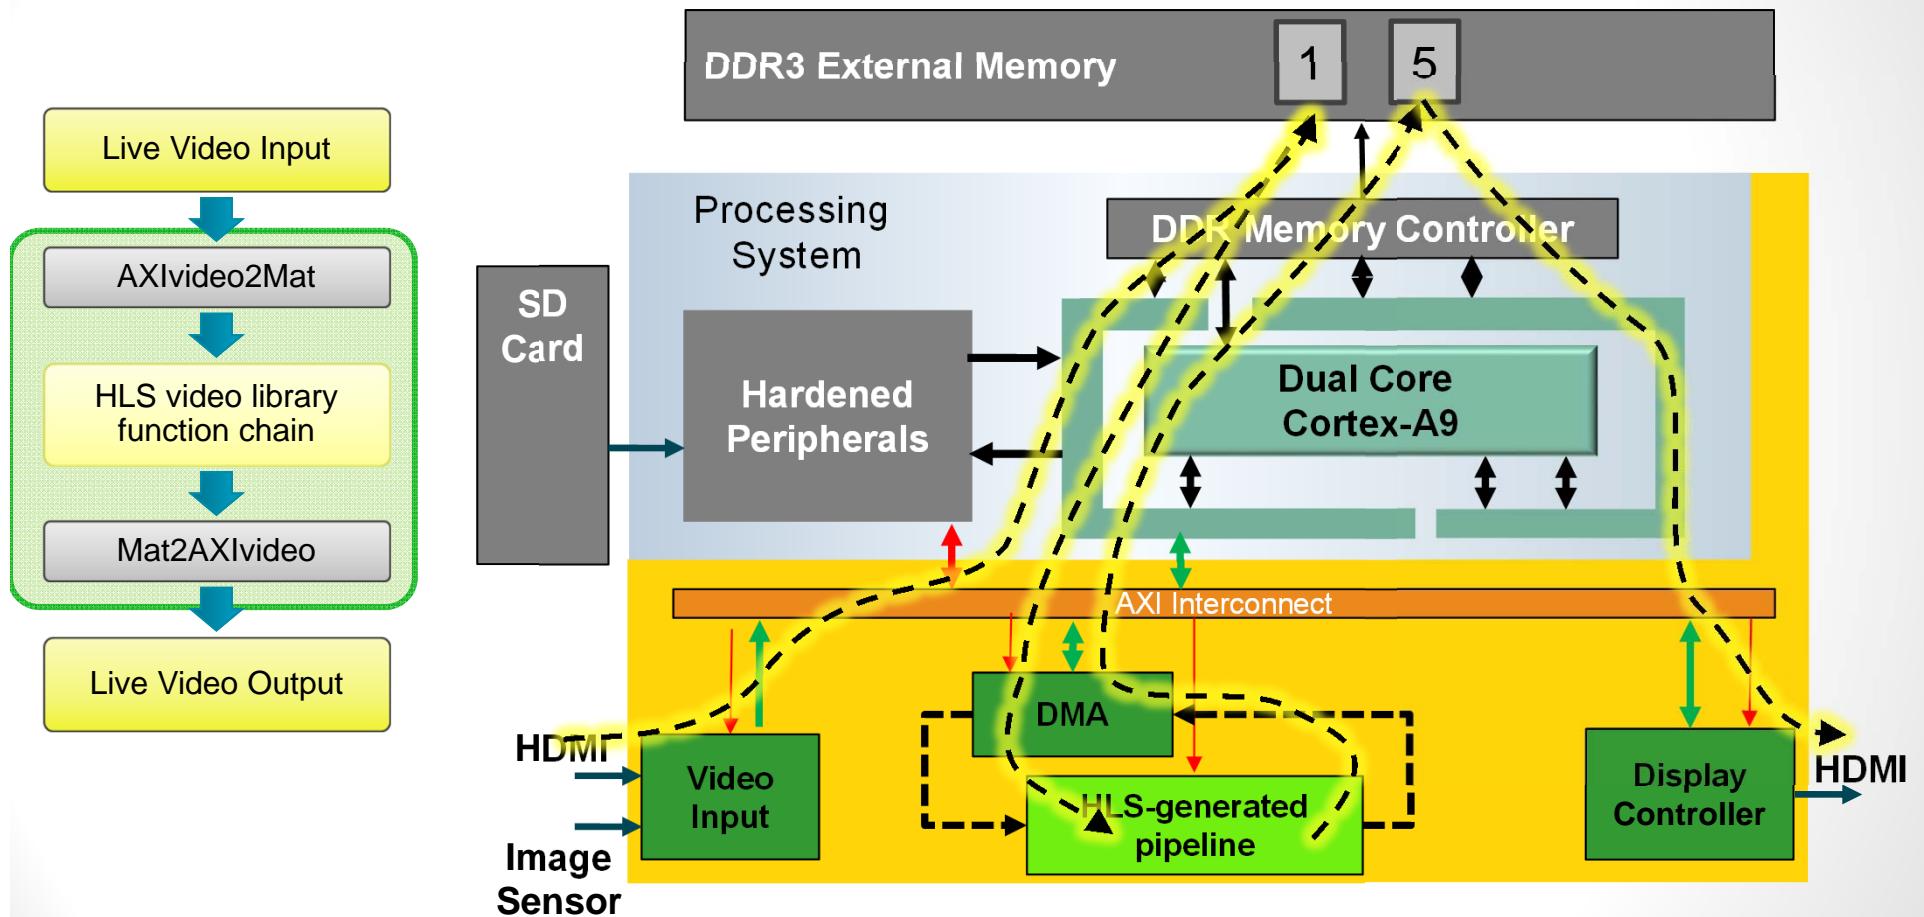


OpenCV Reference



Accelerated OpenCV Application





HLS video libraries

- OpenCV functions are not directly synthesizable with HLS
 - Dynamic buffer allocation
- HLS video libraries are similar to OpenCV image processing functions
 - Require 'streaming' architectures

Video Library functions

- C++ code contained in `hls` namespace
- Similar interface, equivalent behavior with OpenCV, e.g.
 - OpenCV library: `cvScale(src, dst, scale, shift);`
 - HLS video library: `hls::scale<...>(src, dst, scale, shift);`
- Some constructor arguments have corresponding or replacement template parameters: e.g.,
 - OpenCV library: `cv::Mat mat(rows, cols, CV_8UC3);`
 - HLS video library: `hls::Mat<ROWS, COLS, HLS_8UC3> mat(rows, cols);`
- ROWS and COLS specify the maximum size of an image processed

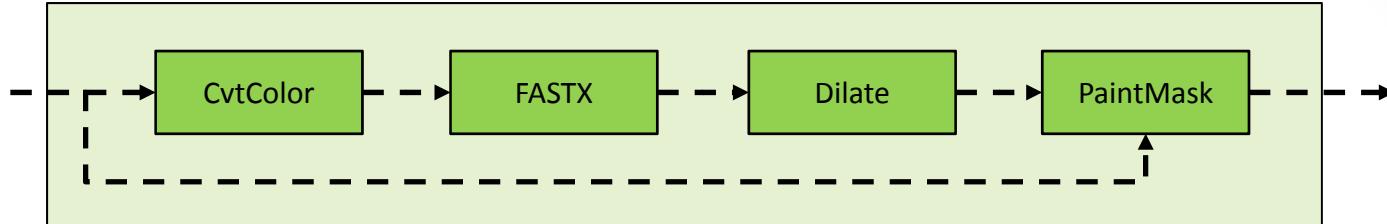
A More Complex Example

- This code is not ‘streaming’ and must be rewritten
 - Random access and in-place operation on ‘dst’

```
void opencv_image_filter(IplImage* img, IplImage* dst ) {  
    IplImage* gray = cvCreateImage(cvSize(img->width,img->height), 8, 1 );  
    cvCvtColor( img, gray, CV_BGR2GRAY );  
    std::vector<cv::KeyPoint> keypoints;  
    cv::Mat gray_mat(gray,0);  
    cv::FAST(gray_mat, keypoints, 20,true );  
    int rect=2;  
    cvCopy(img,dst);  
    for (int i=0; i<keypoints.size(); i++) {  
        cvRectangle(dst,  
                    cvPoint(keypoints[i].pt.x,keypoints[i].pt.y),  
                    cvPoint(keypoints[i].pt.x+rect,keypoints[i].pt.y+rect),  
                    cvScalar(255,0,0),1);  
    }  
    cvReleaseImage( &gray );  
}
```

opencv_top.cpp

Streaming Code



```

void opencv_image_filter(IplImage* src, IplImage* dst)
{
    IplImage* gray = cvCreateImage( cvGetSize(src), 8, 1 );
    IplImage* mask = cvCreateImage( cvGetSize(src), 8, 1 );
    IplImage* dmask = cvCreateImage( cvGetSize(src), 8, 1 );
    std::vector<cv::KeyPoint> keypoints;
    cv::Mat gray_mat(gray,0);

    cvCvtColor(src, gray, CV_BGR2GRAY );
    cv::FAST(gray_mat, keypoints, 20, true);
    GenMask(mask, keypoints);
    cvDilate(mask,dmask);
    cvCopy(src,dst);
    PaintMask(dst,dmask,cvScalar(255,0,0));

    cvReleaseImage( &mask );
    cvReleaseImage( &dmask );
    cvReleaseImage( &gray );
}
  
```

- Synthesizable code
 - Note '#pragma HLS stream'

```

hls::Mat<MAX_HEIGHT,MAX_WIDTH,HLS_8UC3>           _src(rows,cols);
hls::Mat<MAX_HEIGHT,MAX_WIDTH,HLS_8UC3>           _dst(rows,cols);
hls::AXIVideo2Mat(input, _src);
hls::Mat<MAX_HEIGHT,MAX_WIDTH,HLS_8UC3>           src0(rows,cols);
hls::Mat<MAX_HEIGHT,MAX_WIDTH,HLS_8UC3>           src1(rows,cols);
#pragma HLS stream depth=20000 variable=src1.data_stream
    hls::Mat<MAX_HEIGHT,MAX_WIDTH,HLS_8UC1>           mask(rows,cols);
    hls::Mat<MAX_HEIGHT,MAX_WIDTH,HLS_8UC1>           dmask(rows,cols);
    hls::Scalar<3,unsigned char> color(255,0,0);
    hls::Duplicate(_src,src0,src1);
    hls::Mat<MAX_HEIGHT,MAX_WIDTH,HLS_8UC1>           gray(rows,cols);
    hls::CvtColor<HLS_BGR2GRAY>(src0,gray);
    hls::FASTX(gray,mask,20,true);
    hls::Dilate(mask,dmask);
    hls::PaintMask(src1,dmask,_dst,color);
    hls::Mat2AXIVideo(_dst, output);

```

1080P60 Corner Detection



FPGA Pipelines

- Efficiently implement real-time video processing
- Throughput stays constant with more complex programs

Performance of Xilinx HLS video library

	Throughput * (Megapixels per sec)	Acceleration * (vs. OpenCV on ARM)
FAST9 Corner Detection	124	10x
Canny Edge Detection	124	14x
Harris Corner Detection	124	50x
Erosion/Dilation	124	5x
5x5 convolution	124	27x

* FPGA target is 1080P60
compared to single core 1GHz Cortex-A9

Conclusion

- Trends in vision processing
 - Higher resolution and more cameras -> smarter cameras
- Key bottlenecks in embedded video processing
 - Memory + computation -> dedicated processing pipelines
- Using OpenCV in embedded video processing
 - Run on ARM for non-bottleneck portions
 - Migrate to NEON acceleration for 2-20x speedup
 - Use as golden reference for HLS design in FPGAs for 10-100x speedup and real-time processing





Stephen Neuendorffer

Principal Engineer

Xilinx

2100 Logic drive
San Jose, California
USA

Phone: +1 (408) 626-6359

Email: stephenn@xilinx.com

www.xilinx.com

Stereo Camera Architecture

